

Lifecycle OS — Smart Brain

One system that reads our data, watches every competitor, plans the next 15 days, generates every asset (mailers, Google/Meta/TikTok ads, landing pages) and talks to customers in a human voice. Humans verify; the brain executes.

- Live at /brain · DB-linked · daily cron · phase-2 live-push excluded by design
- Built on the existing shell · Supabase · 6-provider LLM cascade · ~\$0 infra
- Now includes the V2 Lifecycle OS layer (2026-07): Mailer Calendar, UK Non-Engagers, Social OS

The marketing team's week, before and after

Task	Manual today	With Smart Brain	Saved
Performance digging	~6 h/wk in sheets	Daily auto-analysis + chat console	~5.5 h
Calendar planning	~4 h/wk, gut-driven	15-day rolling, self-reviewing daily	~3.5 h
Mailer + ad + LP creation	~2 days per set	Full funnel in <60s, on-brand by construction	~85%
Competitor monitoring	Ad-hoc inbox digging	Continuous isolated stream + signals	~2 h
Customer education	Headcount-limited	Voice agents 24x7, every market	scale

Six modules, clean contracts

- 1 · Knowledge Base — creative-level index of what worked
- 2 · Analysis Engine — thresholds, cohorts, daily sync
- 3 · Competitor Stream — isolated; advisory signals only (never qualifies our winners)
- 4 · Calendar Intelligence — 15-day rolling, festival-aware, self-reviewing
- 5 · Generation Engine — mailers, Google/Meta/TikTok ads, landing pages as platform-ready objects
- 6 · Human Review — every campaign verified at launch; hard weekly recalibration gate

Daily is automated. Weekly is human.

- Daily 03:30 UTC cron, no humans: re-extract festivals & peaks from history
- Full re-analysis vs the DB — never follow the calendar blindly, update it
- Apply human feedback, swap weak angles, extend the 15-day horizon
- Fold MVT winners into learned weights + competitor signals
- Generate complete asset sets for approved slots due within 3 days
- HUMAN GATE: approve → final; reject → reasons feed back. Recalibration >7 days old hard-blocks auto-approval

Operate it like ChatGPT, grounded in live data

- Claude-style chat that plans multi-step work and calls real tools
- Collapsible tool trace: every call, its args, and results
- Evidence contract: quote exact tool-sourced figures, name the metric + expected impact, state a hypothesis, quote competitor benchmarks
- Works across the entire 6-provider waterfall including free tiers

A 15-day plan that argues for itself

- Each slot: date × market × cohort (live size) × channel × theme/angle × hero product
- Every entry carries a rationale and a confidence score
- Festival-weighted and send-time tuned; MVT loop records per-variant outcomes
- Daily review narrates what changed and why (angle swaps, skips from your feedback)

One approval → a complete funnel

- Approved slot (cohort + angle + theme) fans out to creative, landing, follow-up
- HTML mailer (brand-locked: palette · Lao MN · banned-phrase scrub)
- Google RSA · Meta · TikTok copy + baked-text creatives
- Two landing pages (for the mailer and for the ads) hosted at /lp/:id
- Audience specs (retargeting 14d, lookalike 1-3%) — stored as platform-ready JSON; phase-2 push needs zero refactor

Telecalling held ROAS 3.5-4. This reproduces it at ~\$0.

- Value justification is conversational: duration, benefits, per-cup math
- Brand / product / collection agents — voice in, voice out, honest by design
- One script tag on Shopify routes to the right specialist; mailers & LPs deep-link to /agent
- Knowledge = catalog slice + scraped official product pages + brand KB

How we know it's right — and how it improves

Surface	Acceptance criteria (launch)	Improvement loop
Analysis	Scores reproducible from raw metrics; pass-rate 20-60%	Weekly recalibration adjusts thresholds
Calendar	≥80% of tentative slots approved without edits by wk 2	Rejections suppress angles; MVT winners boost weights
Assets	100% brand-compliant by code; ≥70% approved w/o edits	Rejection notes fed into next prompt
Voice	0 medical claims · 0 banned phrases · ≥85% resolution	Persona + knowledge updates from transcript audit
System	Daily cron ≥99%; recalibration never >7d	Confidence relaxes review load only on proof

What we expect it to move

- Bandwidth: about -70% production hours per campaign set (team shifts to verification + strategy)
- Send-to-plan latency: 15 days → daily, re-validated every morning
- Creative win-rate: +20-30% (only threshold-passing angles seed new campaigns; MVT compounds)
- Education channel ROAS: ~3.5x via voice agents
- Cost: ~\$0 (free-tier DB + existing Vercel + free LLM tiers); Risk: gated (human-verified, weekly recalibration)

From today to autopilot

- Week 0 (done): Smart Brain live at /brain; cron + review queue + agents ready
- Week 1: point the linked DB at real Shopify/Klaviyo/ads exports; harden RLS; first recalibration
- Week 2-3: shadow mode — brain plans & generates, team ships manually; tune thresholds weekly
- Week 4+: phase 2 — Klaviyo + ad-platform push on stored objects; performance flows back; loop closed

WHERE THIS GOES

VAHDAM first. SaaS later.

- Nothing in the brain is VAHDAM-specific by code — brand kit, thresholds, catalog, agents are DB rows
- Multi-tenancy is one linked DB per brand away
- Wedge: D2C brands with scattered history, no analyst bandwidth, a price-justification problem
- Pricing instinct: per-brand platform fee + per-generated-campaign / per-agent-conversation usage

The current-state layer

- V1/V2 taxonomy adopted (2026-07-03): V1 = legacy base; V2 = Lifecycle OS additions
- V2: cohort Mailer Calendar, UK Non-Engagers hub, tier-routed cascades + video-core, Social Media OS, retention/influencer KB
- Access: the demo/mock gate was removed — every signed-in user gets full live access
- Design: UI locked to a single forest-green theme (dark/dusk/light switcher removed)
- Ops: domain + OAuth migration tooling (migrate-domains / migrate-oauth) keeps Google sign-in working on new domains

LIVE

See it running

- Console: /brain · Operator: /chaigpt · Buyer agent: /agent
- This deck: /deck · PRD: /prd-deck
- Everything on the brand's forest-green canvas, gold + cream type, Lao MN / Proxima Nova